




# Peer Discovery in Ethereum

py-libp2p  
(Robert Zajac, Alex Haynes, et. al)



# Who are we

- Four University of Pennsylvania (UPenn) CS undergrads (4th year)
- py-libp2p team
  - Python client of the libp2p standard



Robert Zajac



Alex Haynes




Aspyn Palatnick



ZX Zhang

# Overview

1. State of peer discovery in Ethereum - **discv4**
  - a. Weaknesses & vulnerabilities
2. Current alternatives
  - a. **libp2p** and its solutions (Kademlia DHT)
3. [Discussion] Peer discovery in Ethereum 2.0
  - a. “discv5”
  - b. Desirable properties & improvements
  - c. Unifying libp2p + Eth 2.0




*Disclaimer:* information not canonical



# Peer Discovery in Ethereum

- Discovery layer of *devp2p* (Ethereum networking stack)
  - Geth, Trinity implement discv4
- Kademlia-inspired DHT
  - Node ID: 64-byte elliptic curve public key
  - Kademlia ID (used in routing table):  $\text{SHA256}(\text{node id})$



# discv4 RPCs (simplified)



1. *Ping*(**from**, **to**)
  - a. Test connection to a peer. **from** + **to** contain ip/port details
2. *Pong*(**to**, **ping-hash**)
  - a. Respond to *Ping*. Clients ignore *Pong* with invalid **ping-hash**.
3. *FindNode*(**id**)
  - a. Ask for the  $k$  ( $k = 16$ ) nodes closest to the node **id**.
4. *Neighbors*(**nodes**)
  - a. Respond to *FindNode*. **nodes** = [(ip, port, node id)...

All packets include *expiration timestamp* to prevent packet replay attacks

# discv4 Details



- Initial peer discovery
  - Boot process queries *bootnodes* for closest neighbors (*FindNode*), repeats on result
- Routing table (Kad-compliant)
  - Consists of *k-buckets* (list of recently seen nodes)
  - For each  $0 \leq i < 256$ , node keeps a *k-bucket* of nodes at distance  $(2^i, 2^{i+1})$ 
    - “Distance” = XOR of Kad IDs
  - *k-buckets* populated upon *FindNode*; only most recent *k* peers kept (time of last-successful *Ping/Pong* exchange)
- Bonding (endpoint proof)
  - Guarantee authenticity of incoming request
    - Avoid DoS attacks, where attacker spoofs *FindNode* from a peer and floods it with *Neighbors* response
  - Peers *Ping* each other and validate response (via hash) before issuing *FindNode*

# discv4 Drawbacks



- Implementation difficulties around bonding; code complexity
- Bonding process is imprecise
  - Uncertain if *Pong* is processed, peer may consider other peer bonded when they are not
- Process for recovering full information about other nodes is not specified by discv4
  - Node IDs are found by looking at the public key used to sign UDP packets after the bonding process
- Multiple attacks (e.g. Eclipse) have been identified on Kademlia





# Peer Discovery Alternatives


# Peer discovery alternatives



- Kademlia DHT
  - Differences between Kademlia and discv4
    - Kademlia offers both finding nodes and storing/finding values (the latter is unused by discv4)
    - Discv4 additionally uses 64-byte public keys as node IDs in addition to having a Kad ID (SHA256(node ID))
- mDNS
  - Mostly used in for small/local networks
- Address Advertisement
- Many protocols implement custom peer discovery in-house (see BitTorrent)
- ...libp2p!

# What is libp2p?


- *Not* devp2p
- Modular p2p networking library
- Original networking stack of IPFS
  - Backed by Protocol Labs
- Clients in Go and JS
  - Python (us), Rust, Java and Haskell coming soon



# libp2p Peer discovery

---

- Modular design of library allows use of any module conforming to a “discovery interface”
  - Kademlia DHT (most common)
  - mDNS
  - “discv5” could easily be plugged in
- libp2p peer discovery solutions openly specified/documented
- py-libp2p developing a Kad DHT solution
  - Current goal: backwards compatibility with devp2p





# Ethereum 2.0 Peer Discovery

# Current work on “discv5”

---

- “discv5” is a misnomer; refers to next generation of peer discovery in Ethereum
- Geth upgraded to a “discv5” implementation as part of [1.5 release](#) (2016)
  - Other clients on hold
- Specs are unofficial, and still in the works (why we’re here)



# Future of peer discovery - *talking points*



## Desirable Qualities of Spec

- Returning ENR Records?
  - *FindNodes* provide node metadata instead of simply address information
- Masking algorithm?
  - Avoid information leak via signed UDP packets
- Alternative to bonding process?
  - More precise, resilient to network

## Integrating libp2p + Eth 2.0

- Leveraging rich libp2p specs + documentation?
- Code reuse between IPFS + Ethereum?
- Joint effort by communities?

**Let's discuss**



**Thank you!**